

# Cluster and Cloud Computing

## Assignment I Report

### Authors

[yuan.gao.2@student.unimelb.edu.au](mailto:yuan.gao.2@student.unimelb.edu.au)

[yzhao5067@student.unimelb.edu.au](mailto:yzhao5067@student.unimelb.edu.au)

### Login and Initialization

1. use **ssh-key** to login without password, copy public key to the remote host and set ssh config file and alias command locally to login by simply command 'spartan'.
2. Create shell scripts `initial_env.sh` and set alias command 'alias inienv='source /home/ygao3631/initial\_env.sh' in ~/.bashrc 'to **conveniently initialize each time**, The function includes module purge and Load spartan & foss/2022a & Python/3.10.4 & Scipy

```
garvyn@Yuans-MacBook-Air: ~ % spartan
Register this system with Red Hat Insights: insights-client --register
Create an account or view all your systems at https://red.ht/insights-dashbo
Last login: Mon Mar 31 22:15:32 2025 from 194.127.105.122

Welcome to Spartan, the general purpose High Performance Computer system.
Use of Spartan is governed by the Research Computing Services policies - htt
Do not run programs or code on the login node. Submit them to the queue with
'sinteractive'. Details on these commands and various other examples of how
can be found in /apps/examples or by typing "man spartan".
To see the usage of the system, type spartan-weather
```

```
[ygao3631@spartan-login3 ~]$ inienv
Module purge : SUCCESS
Module spartan: SUCCESS
Module foss/2022a: SUCCESS
Module Python/3.10.4: SUCCESS
Module SciPy-bundle/2022.05: SUCCESS
All Done !!
```

### The slurm scripts for submitting the job

The final version:

```
#!/bin/bash
#SBATCH --job-name=mastodon_analysis
#SBATCH --output=mastodon_analysis_%j.out
#SBATCH --error=mastodon_analysis_%j.err
#SBATCH --nodes=[node_amount]
#SBATCH --ntasks-per-node=[core_amount]
#SBATCH --time=04:00:00
#SBATCH --mem=[memory per-core]

mpirun -n 8 python mastodon_analysis.py
# or sun -n 8 python mastodon_analysis.py
```

For different Jobs, edit '--nodes='; '--ntasks-per-node=' to adjust the different circumstances.

- **2 nodes ,4 cores for each:** set '--nodes=2'; '--ntasks-per-node=4'.
- **1 node, 8 cores for each:** set '--nodes=1'; '--ntasks-per-node=8'.
- **1 node ,1 core for each:** set '--nodes=1' ; '--ntasks-per-node=1'.

```
1 #!/bin/bash
2 #SBATCH --job-name=mastodon_analysis
3 #SBATCH --output=mastodon_analysis_11%j.out
4 #SBATCH --error=mastodon_analysis_11%j.err
5 #SBATCH --nodes=1
6 #SBATCH --ntasks-per-node=1
7 #SBATCH --time=04:00:00
8 #SBATCH --mem=8G
9
10 mpiexec -n 1 python mastodon_analysis.py
```

```
1 #!/bin/bash
2 #SBATCH --job-name=mastodon_analysis
3 #SBATCH --output=mastodon_analysis_%j.out
4 #SBATCH --error=mastodon_analysis_%j.err
5 #SBATCH --nodes=2
6 #SBATCH --ntasks-per-node=4
7 #SBATCH --time=04:00:00
8 #SBATCH --mem=8G
9
10 mpiexec -n 8 python mastodon_analysis.py
```

```
1 #!/bin/bash
2 #SBATCH --job-name=mastodon_analysis
3 #SBATCH --output=mastodon_analysis_%j.out
4 #SBATCH --error=mastodon_analysis_%j.err
5 #SBATCH --nodes=1
6 #SBATCH --ntasks-per-node=8
7 #SBATCH --time=04:00:00
8 #SBATCH --mem=8G
9
10 mpiexec -n 8 python mastodon_analysis.py
```

## Approach to build and parallelize the code

### *From stream-based to File-slicing-based MPI design*

- **Stream-based MPI design:** The first try uses send and recv to stream data from rank 0 to other processes. While it gives more control over memory and avoids partial line handling, it's slower because rank 0 reads the entire file alone and becomes the bottleneck for both I/O and communication. This approach runs well on 16m data but take much more time than expected on 144G data: The second version is less efficient might because rank 0 handles all the file reading, which creates a severe I/O bottleneck. Other ranks must wait for data to be sent, causing idle time. Additionally, the frequent use of send and recv introduces high communication overhead, especially when processing large files like 144GB. This serialized workflow limits parallelism and slows down the entire program.
- **File-slicing-based MPI design:** The key optimization is replacing centralized streaming with parallel file reading. Each process reads and processes data independently, and only the final results are gathered, making the whole program much faster and more scalable. The second version lets each process read its own part of the file directly using seek, so all ranks process data in parallel. This greatly improves speed by removing the central bottleneck and reducing communication overhead.
- **Optimization:** Replaced comm.send/recv with direct file slicing using f.seek() ; Let each process handle its own I/O and parsing logic ; Only final aggregation is done on Rank 0 using comm.gather().

```
def main():
    comm = MPI.COMM_WORLD
    rank = comm.Get_rank()
    size = comm.Get_size()

    filename = 'large-144G.ndjson'
    chunk_size = 1024 * 1024 * 100 # 100MB

    hour_sentiment = defaultdict(float)
    user_sentiment = defaultdict(float)
```

```

with open(filename, 'r', encoding='utf-8') as f:
    f.seek(0, 2)
    file_size = f.tell()
    chunk_size = file_size // size
    start = rank * chunk_size
    end = start + chunk_size if rank != size - 1 else file_size

    f.seek(start)
    if rank != 0:
        f.readline()

    pos = f.tell()
    while pos < end:
        line = f.readline()
        if not line:
            break
        pos = f.tell()
        created_at, sentiment, user_id, username = parse_line(line)
        if created_at and sentiment is not None and user_id and username:
            try:
                dt = datetime.fromisoformat(created_at.replace('Z', '+00:00'))
                hour = dt.strftime('%Y-%m-%d %H:00')
                hour_sentiment[hour] += sentiment
                user_sentiment[username] += sentiment
            except ValueError:
                continue

```

- **Fault tolerance Design:** The code include fault tolerance to deal with the dirty data(e.g. malformed JSON, missing fields) with out crashing

```

def parse_line(line):
    try:
        data = json.loads(line) # May raise JSONDecodeError
        doc = data.get('doc', {}) # Prevents KeyError
        created_at = doc.get('createdAt', None) # Returns None if missing
        sentiment = doc.get('sentiment', None)
        account = doc.get('account', {})
        user_id = account.get('id', None)
        username = account.get('username', None)
        return created_at, sentiment, user_id, username
    except json.JSONDecodeError:
        return None, None, None, None # Marks invalid records

```

- **Collect Results in rank 0**

```

all_hour_sentiment = comm.gather(hour_sentiment, root=0)
all_user_sentiment = comm.gather(user_sentiment, root=0)


if rank == 0:
    combined_hour = defaultdict(float)
    combined_user = defaultdict(float)

```

```
for hs in all_hour_sentiment:
    for hour, sentiment in hs.items():
        combined_hour[hour] += sentiment

for us in all_user_sentiment:
    for user, sentiment in us.items():
        combined_user[user] += sentiment
```

		1	Rank 1: Processing time: 247.96 seconds				
		2	Rank 3: Processing time: 250.94 seconds				
		3	Rank 0: Processing time: 251.00 seconds				
		4	Rank 2: Processing time: 254.91 seconds	4	Rank 2: Processing time: 252.50 seconds		
		5	Rank 4: Processing time: 261.75 seconds	5	Rank 5: Processing time: 252.56 seconds		
		6	Rank 6: Processing time: 261.88 seconds	6	Rank 6: Processing time: 252.67 seconds		
		7	Rank 7: Processing time: 263.95 seconds	7	Rank 0: Processing time: 254.38 seconds		
		8	Rank 5: Processing time: 266.75 seconds	8	Rank 7: Processing time: 254.97 seconds		
1	Rank 0: Processing time: 2596.89 seconds	9		9			
2		10	5 Happiest Hours:	10	5 Happiest Hours:		
3	5 Happiest Hours:	11	2025-01-01 00:00 with sentiment score 206.15209972304456	11	2025-01-01 00:00 with sentiment score 206.15209972304456		
4	2025-01-01 00:00 with sentiment score 206.15209972304362	12	2024-12-31 23:00 with sentiment score 187.465476632202	12	2024-12-31 23:00 with sentiment score 187.465476632202		
5	2024-12-31 23:00 with sentiment score 187.40547663220213	13	2025-01-01 05:00 with sentiment score 135.92521743446807	13	2025-01-01 05:00 with sentiment score 135.92521743446807		
6	2025-01-01 05:00 with sentiment score 135.92521743446807	14	2024-12-24 22:00 with sentiment score 117.28631361201234	14	2024-12-24 22:00 with sentiment score 117.2863136120121		
7	2024-12-24 22:00 with sentiment score 117.28631361201234	15	2024-12-25 15:00 with sentiment score 114.60195453225161	15	2024-12-25 15:00 with sentiment score 114.60195453225167		
8	2024-12-25 15:00 with sentiment score 114.60195453225161	16		16			
9		17	5 Saddest Hours:	17	5 Saddest Hours:		
10	5 Saddest Hours:	18	2024-11-06 07:00 with sentiment score -373.7672479738356	18	2024-11-06 07:00 with sentiment score -373.76724797383497		
11	2024-11-06 07:00 with sentiment score -373.7672479738356	19	2024-09-11 01:00 with sentiment score -305.57606693741	19	2024-09-11 01:00 with sentiment score -305.5760669374102		
12	2024-09-11 01:00 with sentiment score -305.57606693741	20	2025-01-30 17:00 with sentiment score -256.5684250419835	20	2025-01-30 17:00 with sentiment score -256.5684250419817		
13	2025-01-30 17:00 with sentiment score -256.5684250419835	21	2025-01-30 18:00 with sentiment score -226.90971970072636	21	2025-01-30 18:00 with sentiment score -226.90971970072644		
14	2025-01-30 18:00 with sentiment score -226.90971970072636	22	2025-02-03 16:00 with sentiment score -223.8334876843535	22	2025-02-03 16:00 with sentiment score -223.83348768435087		
15	2025-02-03 16:00 with sentiment score -223.8334876843535	23		23			
16		24	5 Happiest Users:	24	5 Happiest Users:		
17	5 Happiest Users:	25	gameoflife with sentiment score 9105.741942204384	25	gameoflife with sentiment score 9105.741942204433		
18	gameoflife with sentiment score 9105.741942204384	26	EmojiAquarium with sentiment score 2605.0192646105224	26	EmojiAquarium with sentiment score 2605.0192646105183		
19	EmojiAquarium with sentiment score 2605.0192646105224	27	TheFigen_ with sentiment score 2541.4742243891897	27	TheFigen_ with sentiment score 2541.4742243891906		
20	TheFigen_ with sentiment score 2541.4742243891897	28	choochoo with sentiment score 1978.7847677692084	28	choochoo with sentiment score 1978.7847677692096		
21	choochoo with sentiment score 1978.7847677692084	29	hnb0t with sentiment score 1914.9517912924869	29	hnb0t with sentiment score 1914.9517912927743		
22	hnb0t with sentiment score 1914.9517912924869	30		30			
23		31	5 Saddest Users:	31	5 Saddest Users:		
24	5 Saddest Users:	32	realTuckFrumper with sentiment score -9093.795561168281	32	realTuckFrumper with sentiment score -9093.795561168006		
25	realTuckFrumper with sentiment score -9093.795561168281	33	uavideos with sentiment score -5901.960347558229	33	uavideos with sentiment score -5901.960347558052		
26	uavideos with sentiment score -5901.960347558229	34	TheHindu with sentiment score -5710.722600655962	34	TheHindu with sentiment score -5710.722600655697		
27	TheHindu with sentiment score -5710.722600655962	35	utisbot with sentiment score -4066.1381701402247	35	utisbot with sentiment score -4066.138170138108		
28	utisbot with sentiment score -4066.1381701402247	36	MissingYou with sentiment score -3341.603338119594	36	MissingYou with sentiment score -3341.603338119541		
29	MissingYou with sentiment score -3341.603338119594	37		37			
30		38	Data aggregation time: 0.37 seconds	38	Data aggregation time: 0.78 seconds		
31	Data aggregation time: 0.37 seconds	39	Total execution time: 2597.42 seconds	39	Total execution time: 256.00 seconds		
32	Total execution time: 2597.42 seconds						



## Amdahl'S Law Application

Observed speedup  $\approx 2597 / 256 \approx 10.14x$  ;  $2597 / 267 \approx 9.72x$  **then** Using Amdahl's Law to estimate parallel portion P. One node eight cores & Two nodes four cores:

$$10.14 = 1 / ((1 - p) + (p/8)) \rightarrow P \approx 0.995 \quad (1)$$

for 2 nodes, Still benefits from high parallelism (~99.5%), but performance loss comes from cross-node communication overhead and shared file system I/O contention.

Mode	Time	Speed Up	Performance
1node,1core	2597	1x	No parallelism, CPU and I/O are bottlenecks
1node,8cores	256	10.14x	Near-optimal scaling, minimal communication
2node,4cores	267	9.7x	Effective parallelism, slight inter-node cost

**! Tips** -- More information about CI/CD process can be seen our Github website [https://github.com/GarvynY/Cluster\\_parallel\\_computing\\_spartan.git](https://github.com/GarvynY/Cluster_parallel_computing_spartan.git)